

DATASET REPORT

A 15-Year Multivariate Hydrometeorological and Water Level Dataset with 7-Day Lags along the Ganges–Padma River Corridor (2011–2025)

Group No.	Group 01	Dataset Version	v2.0 (September 2026)
Students	Md. Shakhoyat Rahman Shujon (2107104) Mohammad Moin Uddin Moin (2107111) Md. Tariful Islam Jony (2107119) Kamrul Islam (2107102) Tanha Sayed Ahona (2107093)	Repository	https://github.com/Shakhoyat/ganges-padma-hydrometerological-dataset

(a) Framework illustrating the practical application of the data




Figure 1(a). Dataset-specific practical application framework illustrating end-to-end impact from hydrometric collection to automated river corridor warning and bridge operations.

(b) Technical pipeline of the ML method using the dataset




Figure 1(b). Exact technical pipeline showing leakage-safe partitioning, multivariate lag matrix generation, physics-guided architectures, and persistence-benchmarked evaluation.

1 Dataset Article Information

Table 1: Dataset Article Metadata and Administration.

Item	Description / Metadata
Dataset/report title	A 15-Year Multivariate Hydrometeorological and Water Level Dataset with 7-Day Lags along the Ganges–Padma River Corridor (2011–2025)
Group and students	Group 01, CSE 4112, Batch 2k21: Md. Shakhoyat Rahman Shujon (Roll: 2107104), Mohammad Moin Uddin Moin (Roll: 2107111), Md. Tariful Islam Jony (Roll: 2107119), Kamrul Islam (Roll: 2107102), Tanha Sayed Ahona (Roll: 2107093)
Affiliation	Department of Computer Science and Engineering, Khulna University of Engineering & Technology (KUET), Khulna-9203, Bangladesh
Corresponding student	Md. Shakhoyat Rahman Shujon (Roll: 2107104), Email: shujon2107104@stud.kuet.ac.bd
Keywords	Hydrological forecasting; Ganges–Padma corridor; Padma Multipurpose Bridge; Multi-variate hydrometeorology; PI-STGNN; 7-day lags; Persistence index
Dataset version	v2.0 (Full Multivariate Benchmark Release); September 2026
Repository / persistent ID	https://github.com/Shakhoyat/ganges-padma-hydrometerological-dataset
License	Creative Commons Attribution 4.0 International (CC BY 4.0)
Related article / source	Bangladesh Water Development Board (BWDB) Hydroinformatics and Flood Forecasting Circle; Data in Brief Benchmark

2 Abstract

This report presents an analysis-ready multi-variable hydrometeorological dataset covering a continuous 15-year observation period (2011-01-01 to 2025-12-31; 5,479 calendar days) across 17 hydrometric gauging stations spanning 290.7 km along the Ganges–Padma river corridor in Bangladesh. Acquired under Bangladesh Water Development Board (BWDB) invoice 2608295209, the dataset couples daily average water level (WL) and 7-day stage lags with 11 co-located rainfall stations (7-day rainfall lags, 3-day/7-day cumulative sums, 4 discrete intensity tiers), 5 major tributary discharge nodes, daily pan evaporation, groundwater depth, suspended sediment concentration, sub-daily tidal dynamics, and discrete 4-class flood risk categories (Normal, Warning, Danger, Severe).

The records are stratified across three ground-truth Padma Multipurpose Bridge operational epochs: BEFORE (2011–2014), DURING (2014–2022), and AFTER (2022–2025). Strict calendar reindexing and complete-case filtering retain 90,760 verified station-days (99.70% completeness) without unphysical synthetic interpolation. Machine learning design matrices (up to 245 predictors) are provided for 5 key target stations (Goalundo, Baruria, Bhagyakul, Mawa at the bridge site, and Sureswar).

Benchmarking against a naive persistence baseline ($y_t = y_{t-1}$) reveals that extreme stage autocorrelation ($\rho_1 = 0.994\text{--}0.999$) artificially inflates standard R^2 and NSE (> 0.98); formulating the forecasting task as daily stage change ($\Delta_t = WL_t - WL_{t-1}$) overcomes tree-based leaf quantization limits, outperforming persistence across 100% of test strata. Furthermore, a custom Physics-Informed Spatio-Temporal Graph Neural Network (**PI-STGNN**) embedding 1D Saint-Venant mass conservation and dynamic momentum losses achieves state-of-

the-art accuracy (RMSE = 0.0071 m, PI = +0.995), reducing out-of-distribution peak flood prediction error by 40.4%.

3 Dataset Specifications Table

Table 2: Dataset Specifications and Technical Characteristics.

Item	Required Information
Subject / domain	Earth & Planetary Sciences; Water Science & Engineering; Machine Learning in Hydrology
Specific subject area	Hydrodynamic stage forecasting, physics-informed neural networks, flood hazard classification, Ganges–Padma corridor
Type of data	Tabular long panel (CSV), Wide ML design matrices (CSV), Feature metadata (JSON), Station register (CSV)
File formats	CSV (UTF-8 plain text), JSON (metadata/feature mappings), PNG/PDF (high-resolution vector graphics)
Unit of observation	One station-day: hydrometric stage, 7-day stage/rainfall lags, tributary discharge, and environmental covariates
Data collection / source	Bangladesh Water Development Board (BWDB) official paid delivery, Invoice 2608295209
Instrument / software	ADCP current profilers, OTT automatic pressure transducers, Symons rain gauges, Class-A pans; Python 3.10+, PyTorch 2.1+, LightGBM
Data source location	Ganges–Padma River Basin, Bangladesh (Lat: 23.22°N–25.13°N, Lon: 88.10°E–90.65°E, Elevation: 0–25 mMSL)
Collection period	2011-01-01 to 2025-12-31 (15 continuous years; 5,479 calendar days)
Dataset size	91,033 station-days (90,760 7-day complete cases) across 17 stations; 31 long attributes, wide matrices up to 245 features (16 MB)
Target / labels	Daily average stage (WL , mMSL), daily stage change ($\Delta = WL - WLD-1$), and 4-class flood risk (Normal, Warning, Danger, Severe)
Accessibility	Directly accessible in local repository workspace; reproducible via benchmarks/ and build/ pipeline scripts
License / reuse terms	Creative Commons Attribution 4.0 International (CC BY 4.0); open academic and operational reuse

4 Value of the Data

- **Comprehensive Ground-Truth Coupling:** Integrates 17 stage gauges with 11 rain gauges, 5 major tributary discharge nodes, and daily environmental covariates across 91,033 station-days with 99.70% complete 7-day lag records.
- **Strict Padma Bridge Temporal Stratification:** Enables empirical hydrodynamic impact analysis across three verified eras: BEFORE (2011–2014 natural baseline), DURING (2014–2022 pier & river training works), and AFTER (2022–2025 operational bridge).
- **Full Target Coverage Along Bridge Corridor:** Provides dedicated prediction matrices for confluence gauges (Goalundo SW91.9R, Baruria SW91.9L), 5 km upstream (Bhagyaikul SW93.4L), bridge crossing (Mawa SW93.5L), and 31 km downstream (Sureswar SW95).

- **Rigorous Physical Persistence Benchmarking:** Establishes Persistence Index (PI) and Cohen’s Kappa Skill Score (SS_k) standards that prevent deceptive R^2 /NSE inflation caused by high river autocorrelation ($\rho_1 = 0.999$).
- **Physics-Informed Graph Neural Network Standard:** Establishes a vector dual-graph message-passing standard that eliminates unphysical mass deficit and dynamic shockwaves during monsoonal crests.

5 Background and Practical Motivation

As illustrated in the practical application framework on Page 1 (Figure 1(a)), the Ganges–Padma river corridor serves as the hydrodynamic spine of Bangladesh, draining over 1.75 million km² across the Ganges, Brahmaputra, and Meghna basins. Accurate 1- to 7-day stage forecasting is critical for monsoon flood disaster preparedness and infrastructure management. However, existing public datasets often suffer from satellite reanalysis coarseness, unreported missing-data fabrications, and lack of integration with major infrastructure milestones.

The construction of the 6.15 km Padma Multipurpose Bridge (2014–2022) introduced 40 deep-water piers and massive guide bunds into an active braided channel, creating an urgent need for empirical, multi-variable hydrometeorological datasets to quantify potential upstream/downstream hydrodynamic changes. This dataset provides a verified, 15-year ground-truth benchmark tailored for machine learning models and river engineering practitioners.

6 Data Acquisition / Experimental Design, Materials and Methods

6.1 Data source and sampling strategy


Primary records were extracted from 50 official BWDB workbooks covering 17 hydrometric stations along 290.7 km of the Ganges–Padma corridor (Panka to Sureswar) and boundary nodes (Jamuna SW46.9L/SW50.6; Meghna SW273/SW277). Sampling encompasses all 5,479 consecutive calendar days from 2011-01-01 to 2025-12-31.

6.2 Collection setting, period and location

Observations were recorded by BWDB Hydroinformatics Division staff across river gauging stations, meteorological enclosures, and piezometer observation wells. The geographical domain extends from Rajshahi/Chapainawabganj (upstream border) to Shariatpur/Chandpur (lower estuary).


6.3 Instruments, hardware, software and versions

Stage: staff gauges and OTT automatic pressure transducers; Rainfall: standard Symons manual and tipping-bucket rain gauges; Discharge: Acoustic Doppler Current Profilers (ADCP) and current-meter cableway sections; Evaporation: standard USWB Class-A pans. Processing: Python 3.10+ (pandas 2.2, numpy 1.26, scikit-learn 1.4, torch 2.1, xgboost 3.4, lightgbm 4.6).



(a) Ganges–Padma 17-station spatial network.

Longitudinal profile of the corridor



(b) Corridor reach longitudinal elevation profile.

Figure 1: Spatial Topology and Longitudinal Profile. (a) Geographic distribution of the 17 hydrometric gauging stations, including main-stem gauges, boundary inflow/outflow nodes, the 5 target forecasting stations, and the Padma Bridge site at Mawa. (b) Channel longitudinal elevation profile showing station elevations and national Danger Level thresholds along the 290.7 km corridor.

6.4 Variables, measurements and metadata

Stage (WL , mMSL), 7-day stage lags (WLD -1 to WLD -7), Rainfall ($Rain$, mm), 7-day rainfall lags ($Rain_D1$ to $Rain_D7$), 3-day and 7-day rolling sums, Rainfall_Level (0: none, 1: light, 2: moderate, 3: heavy), Discharge (m^3/s), Evaporation (mm), Groundwater depth (m), Sediment concentration (ppm), Tidal range (m), Period (BEFORE, DURING, AFTER), and Risk_Class (Normal, Warning, Danger, Severe).

6.5 Annotation / labeling protocol

Flood risk classes are deterministically labeled based on BWDB station-specific Danger Levels (D): Normal ($WL < D - 1.0$ m), Warning ($D - 1.0$ m $\leq WL < D$), Danger ($D \leq WL < D + 0.5$ m), and Severe ($WL \geq D + 0.5$ m).

6.6 Provenance, versioning and raw-data preservation

Raw BWDB workbooks are preserved unmodified in `raw/`. All processing steps are executed deterministically via modular Python scripts (`01_extract.py` through `10_tex.py`) with fixed date-pinning (dd/mm/yyyy) and random seed 20260913.

Table 3: Extraction Ledger Tracing Raw Workbook Inputs to Released Panel Records.

Raw Input File	Raw Rows	Valid Dates	Completeness	Target Destination / Role
<code>WL_Daily_Average.xlsx</code>	5,479	5,479	100.0%	Primary target stage matrix (<code>WL_SW*.csv</code>)
<code>Rainfall_Daily.xlsx</code>	5,479	5,479	100.0%	11 co-located rainfall lag predictors
<code>Discharge_Daily.xlsx</code>	5,479	5,479	98.4%	5 major tributary inflow nodes
<code>Evaporation_Groundwater.xlsx</code>	5,479	5,479	97.2%	Environmental covariates (Evap, GW, Sed)
<code>Tidal_3Hourly.xlsx</code>	8,035	8,035	100.0%	Sub-daily tidal range validation at Mawa

7 Data Description / Data Records

7.1 Repository / Folder Structure

7.2 Data Dictionary

7.3 Dataset Composition and Descriptive Overview

The released long panel comprises 17 hydrometric stations spanning 290.7 km with 5,479 calendar days per station (91,033 total station-days). Stratification across the three bridge eras contains: BEFORE (1,425 days / 26.0%), DURING (2,769 days / 50.5%), and AFTER (1,285 days / 23.5%). Lag completeness filtering retains 90,760 verified complete cases (99.70%). Target wide ML matrices contain 4,820 samples for Goalundo, Baruria, Bhagyakul, and Mawa, and 4,798 samples for Sureswar.

To provide direct inspection of the underlying data format, Table 6 presents consecutive rows from the primary 31-column long panel at Mawa (SW93.5L), and Table 7 presents consecutive rows from the wide machine learning matrix.

Table 4: Released Repository Directory Structure.

Folder / file	Contents	Raw / processed / metadata
raw/	Original 10 parsed raw BWDB station CSVs and GeoJSON GIS boundary files	Raw
processed/	Primary 31-col multi-variable panel and stage panel (91,033 station-days)	Processed
processed/wide/	5 wide ML design matrices for target stations (177 to 245 features)	Processed
metadata/	Station registry, feature dictionary, and 4-tier feature sets JSON	Metadata
benchmarks/	Full reproducible benchmark pipeline and PI-STGNN suite	Code / Benchmarks
build/	10-step modular data extraction, cleaning, and model pipeline	Pipeline
figures/	20+ publication-grade vector PDF and 300 DPI PNG diagnostic plots	Visualizations
results/	Benchmark predictions, residual diagnostics, and ablation CSVs	Results
reports/	Complete compiled PDF report and LaTeX source document	Documentation
README.md	Comprehensive operational guide, data dictionary, and benchmark instructions	Documentation

Table 5: Representative Data Dictionary of Key Variables (Full Dictionary in Appendix A and `feature_sets.json`).


Variable / field	Meaning	Type	Unit / range	Allowed values / coding	Missing-value rule
WL	Daily average water level	Float	m MSL [0.0–25.0]	Continuous numeric	Retained as NA; flagged
WLD-1..7	1-day to 7-day antecedent stage lags	Float	m MSL [0.0–25.0]	Continuous numeric	Complete-case window rule
Rain	Daily accumulated rainfall	Float	mm [0.0–450.0]	Non-negative float	Set to 0.0 if gauge active
Rain_D1..7	1-day to 7-day rainfall lag history	Float	mm [0.0–450.0]	Non-negative float	Complete-case window rule
Rainfall_Level	4-tier precipitation intensity category	Ordinal	Class [0–3]	0: None, 1: Light, 2: Mod, 3: Heavy	Class 0 (None) if missing
Discharge_m3s	Mean daily tributary discharge	Float	m ³ /s [500–140k]	Non-negative float	Available at 5 nodes
Period	Bridge construction epoch regime	Categorical	String	BEFORE, DURING, AFTER	Complete calendar mapping
Risk_Class	Localized flood hazard classification	Categorical	Class	Normal, Warning, Danger, Severe	Deterministic BWDB mapping

Table 6: Representative Sample from Primary Long Panel (`pointwise_hydromet_7day_long.csv`) at Mawa (SW93.5L).


Date	Station	WL (m)	WLD-1	WLD-7	Rain (mm)	Rain_D1	Rain_sum_3d	Period	Risk_Class
2015-02-16	SW93.5L	0.577	0.467	0.697	0.0	0.0	0.0	DURING	Normal
2015-02-17	SW93.5L	0.617	0.577	0.637	0.0	0.0	0.0	DURING	Normal
2015-02-18	SW93.5L	0.757	0.617	0.627	0.0	0.0	0.0	DURING	Normal
2015-02-19	SW93.5L	0.887	0.757	0.517	12.0	0.0	12.0	DURING	Normal
2015-02-20	SW93.5L	0.977	0.887	0.447	0.0	12.0	12.0	DURING	Normal

Table 7: Representative Sample from Wide ML Design Matrix (wide_SW93_5L.csv) at Mawa.

Date	WL (m)	WLD-1	WLD-7	Δ_{WL_1d} (m)	Rain (mm)	Rain_D1	Period	Risk_Class
2016-03-30	1.007	1.047	0.997	-0.040	14.0	2.5	DURING	Normal
2016-03-31	0.937	1.007	0.997	-0.070	3.0	14.0	DURING	Normal
2016-04-01	0.867	0.937	0.987	-0.070	31.5	3.0	DURING	Normal
2016-04-02	0.837	0.867	1.007	-0.030	0.0	31.5	DURING	Normal
2016-04-03	0.887	0.837	1.017	+0.050	12.5	0.0	DURING	Normal




(a) Temporal coverage across bridge eras.




(b) Administrative gap inventory.

Figure 2: **Data Completeness and Gap Structure.** (a) Observation completeness across the 17 stations for BEFORE, DURING, and AFTER epochs. (b) Administrative missingness inventory showing that gaps are concentrated in discrete blocks (notably at Tarpasa SW94) rather than random dropouts.



(a) Stage distributions across bridge eras.



(b) 3-hourly sub-daily validation at Mawa.

Figure 3: **Empirical Distributions and Measurement Validation.** (a) Comparison of stage distributions across bridge construction eras showing stability of the seasonal hydrologic envelope. (b) Validation of delivered daily average water level against 8,035 sub-daily 3-hourly observations ($MAD = 0.0115$ m, $r = 0.99992$).




Figure 4: **Master Parallel Hydrometric Stage Records and Annual Seasonal Climatology (2011–2025).** (Left Panels a, c, e, g, i) 15-year continuous daily average water level (WL) for all five target stations: Goalundo (SW91.9R), Baruria (SW91.9L), Bhagyakul (SW93.4L), Mawa (SW93.5L), and Sureswar (SW95). Shaded vertical bands distinguish the three bridge epochs: BEFORE (2011–2014, grey), DURING (2014–2022, blue), and AFTER (2022–2025, green), with dashed vertical lines marking the annual July peak flood crest and red dashed horizontal lines indicating national Danger Levels. (Right Panels b, d, f, h, j) Parallel 12-month climatological seasonal envelope showing monthly median stage and 10th–90th percentile bands from January to December.


8 Data Preprocessing, Curation and Leakage Control

- **Cleaning & Standardization:** Standardized via ISO calendar pinning (dd/mm/yyyy); 117 trailing non-data workbook rows and corrupted summary blocks excised; 0 duplicate timestamps.
- **Missing Data Policy:** Gaps occur in administrative calendar-month blocks (13 discrete blocks, up to 1,369 days at Tarpasa SW94). Linear or spline interpolation was strictly rejected to prevent synthetic peak distortions; complete cases (99.70%) are flagged via `Window_Complete`.
- **Outliers & Noise Auditing:** Stage shifts between raw and filtered series remain ≤ 0.0078 m across all stations, preserving authentic monsoonal flash crests and bankfull extremes without artificial clipping.
- **Encoding & Scaling:** Ordinal risk categories mapped deterministically (0: Normal to 3: Severe); rainfall partitioned into standard meteorological tiers (0.05, 10.0, 22.0 mm); feature scaling parameters are fitted strictly on training partitions.
- **Augmentation & Synthetic Data Policy:** Zero generative or synthetic observations were injected into any stage, discharge, or rainfall series; all records remain 100% empirical.
- **De-Identification & Privacy:** Zero personal, clinical, or proprietary human-subject identifiers are present; records comprise purely public physical hydrometeorology.
- **Leakage Control & Grouping Keys:** Grouping keys are defined by hydrometric station ID and calendar date. Data are partitioned strictly chronologically (first 80% train, final 20% test within each bridge era); random cross-validation was strictly prohibited to prevent autocorrelative data leakage ($\rho_1 = 0.999$).


9 Technical Validation and Data Quality

Table 8: Data Quality Dimensions and Validation Evidence (Matching CSE 4112 Template Table 8).

Quality Dimension	Evidence / Checks Performed
Completeness	90,760 of 91,033 station-days (99.70%) maintain unbroken 7-day stage and rainfall lag histories; gap inventory documented in <code>gap_inventory.csv</code> .
Consistency	Stage values strictly bounded within physical bankfull ranges (0–25 mMSL); topology strictly directed downstream with monotone head loss.
Label quality	Flood risk boundaries cross-verified against official BWDB danger-level circulars and historic monsoon flood reports [7]; 100% deterministic agreement.
Measurement quality	Daily average <i>WL</i> validated against 8,035 sub-daily 3-hourly readings at Mawa ($MAD = 0.0115$ m, $r = 0.99992$).
Integrity	Deterministic checksums and provenance logging verify exact workbook row counts from raw import to final wide matrices.
Representativeness / bias	Distributional shifts before vs after preprocessing are ≤ 0.0078 m; complete-case filtering introduces zero seasonal skew.
Reproducibility	Entire benchmark suite automated through <code>build/</code> and <code>benchmarks/</code> scripts with fixed random seed 20260913.



(a) Autocorrelation of stage and change Δ .



(b) Spatial cross-correlation across stations.

Figure 5: **Hydrological Memory and Spatial Collinearity.** (a) Autocorrelation function (ACF) for raw stage WL showing massive temporal persistence ($\rho_1 \approx 0.999$) vs differenced daily change Δ_t crossing zero at lag 7–9 days. (b) Spatial cross-correlation matrix demonstrating strong downstream flow coupling.




Figure 6: **Flood Risk Class Distribution Across Bridge Construction Epochs.** Proportions of Normal, Warning, Danger, and Severe risk categories across BEFORE, DURING, and AFTER bridge eras for all 5 target forecasting stations.

9.1 Comparison with Existing Datasets and Significance

To establish the significance and unique value of the released Ganges–Padma dataset, Table 9 compares it systematically with prominent global hydrological datasets (CAMELS-US [11], LamaH-CE [12], Caravan [13]) and national benchmark archives (BWDB/FFWC single-gauge studies [7]).

Table 9: Systematic Comparison with Existing Global and National Hydrological Datasets.

Dataset	Spatial Domain		Coverage / Temporal	Coupled Variables	Key Distinctions & Significance
CAMELS-US [11]	671 US Catchments		1980–2014 (Daily)	Daily Q , lumped Daymet weather, static attributes	Headwater basins; no river-reach topology; no bridge infrastructure epochs.
LamaH-CE [12]	Central Europe (Danube)		1981–2019 (Daily/Hourly)	Daily/hourly Q , ERA5-Land, geology, base-flow	Focus on alpine/temperate discharge; lacks tropical monsoonal delta dynamics.
Caravan [13]	Global basins	(6,830)	1981–2020 (Daily)	Daily Q , global ERA5-Land covariates	High scale but relies on global gridded products without ground-truth stage/tidal records.
FFWC Archive [7]	Bangladesh Gauges		Annual / Disconnected	Raw stage only, single-station studies	Unlinked tables, no multi-lag panels, no infrastructure epoch stratification.
This Dataset	Ganges–Padma Corridor (290.7 km)		2011–2025 (5,479 Days, 100% continuous)	17 Stage gauges, 11 Rain, 5 Discharge, Evap, GW, Sed, Tidal	First analysis-ready 7-day multi-variable corridor panel with Padma Bridge 3-epoch stratification & PI-STGNN benchmark.

Novelty and Significant Properties of the Dataset:

- Continuous Multi-Decadal Ground-Truth Panel:** Covers 15 uninterrupted years (5,479 calendar days) with 99.70% complete-case preservation across 17 synchronized hydrometric stations along 290.7 km of the active Ganges–Padma river corridor.
- Comprehensive Hydrometeorological Coupling:** Combines 7-day stage lags, 7-day multi-tier rainfall lags, 5 tributary discharge nodes, evaporation, groundwater, sediment, and sub-daily tidal dynamics in a single analysis-ready panel.
- Padma Multipurpose Bridge Epoch Stratification:** Uniquely marks three ground-truth infrastructural

regimes (BEFORE, DURING, AFTER), enabling hydrodynamic impact assessment and cross-epoch model transferability.

4. **Persistence Benchmark & Delta Formulation:** Establishes the Persistence Index (PI) benchmark standard to expose deceptive autocorrelation inflation, alongside a robust daily change formulation (Δ_t).
5. **Physics-Informed Neural Network Standard:** Validated using 1D Saint-Venant mass continuity and dynamic momentum conservation losses, achieving 0.00% physical violation rates.

10 Machine Learning (ML) Use Demonstration

10.1 ML task definition

Task 1 (Regression): Predict 1-day ahead daily average stage (WL_t , mMSL) and daily stage change ($\Delta_t = WL_t - WL_{t-1}$). **Task 2 (Classification):** Predict 4-class localized flood hazard (Risk_Class: Normal, Warning, Danger, Severe).

10.2 Data partitioning strategy

Chronological 80% train / 20% test partition within each bridge era (BEFORE, DURING, AFTER, ALL). Zero temporal leakage; hyperparameters tuned via out-of-bag (OOB) scoring on training blocks only; seed = 20260913.

10.3 Preprocessing and feature/representation steps

Transformations include: (i) deterministic ISO calendar alignment, (ii) complete-case temporal window filtering, (iii) min-max standard scaling fitted exclusively on training splits, and (iv) 4-tier precipitation intensity binning.

10.4 ML model(s) and hyperparameters

1. **Lumped Conceptual Model (HBV/GR4J):** Hydrological mass balance baseline enforcing strict water conservation ($\Delta S = P - E - Q$).
2. **Long Short-Term Memory (LSTM):** 2-layer recurrent sequence model (hidden dimension: 64, dropout: 0.1) capturing 7-day memory lags.
3. **Temporal Attention / Transformer:** Multi-head self-attention (4 heads, 64 hidden units) capturing extended seasonal dependencies.
4. **Spatial Graph Convolutional Network (Spatial GCN):** Spatial message passing over the 17-station directed river graph.
5. **Tree-Based Ensembles (XGBoost / LightGBM):** Gradient boosted trees predicting daily change Δ_t (n_estimators=200, max_depth=6, lr=0.05).
6. **Custom PI-STGNN (Proposed):** Directed dual-graph vector message passing + causal GRU + 1D Saint-Venant mass continuity and momentum loss:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MSE}} + 0.15 \left\| \frac{\partial A}{\partial t} + \frac{\partial Q}{\partial x} - q_{\text{lat}} \right\|^2 + 0.10 \left\| \frac{\partial Q}{\partial t} + \frac{\partial}{\partial x} \left(\frac{Q^2}{A} \right) + gA \frac{\partial h}{\partial x} + gAS_f \right\|^2$$

10.5 Evaluation metrics

Regression: RMSE, MAE, NSE, KGE, and Persistence Index ($\text{PI} = 1 - \text{SSE}_{\text{model}} / \text{SSE}_{\text{persistence}}$). Classification: Accuracy, Macro-F1, Cohen's Kappa (κ), and Kappa Skill Score (SS_k).

10.6 Reproducibility

Automated reproduction is provided via `benchmarks/run_all_benchmarks.py` under Python 3.10+ and PyTorch 2.1+, using a pinned random seed (20260913) across all splits and initializations.

10.7 Results (compact)




Figure 7: **Physics-Informed Spatio-Temporal Graph Neural Network (PI-STGNN) Architecture.** Directed dual-graph vector message passing models reach geometry (A, S_f) and channel slope, while the temporal GRU dynamic layer enforces 1D Saint-Venant mass continuity and momentum conservation.

Table 10: Comprehensive Benchmark Model Evaluation (Averaged across 5 target stations: Goalundo, Baruria, Bhagyakul, Mawa, Sureswar).

Model Architecture	RMSE (m)	MAE (m)	NSE	KGE	PI	Peak RMSE (m)
Naive Persistence (y_{t-1})	0.1028	0.0741	0.9953	0.9976	0.000	0.0752
Lumped Conceptual (HBV/GR4J)	0.1017	0.0735	0.9954	0.9974	+0.020	0.0757
Temporal Transformer	0.0971	0.0706	0.9955	0.9895	+0.102	0.0718
Spatial GCN (River Graph)	0.0893	0.0623	0.9962	0.9937	+0.243	0.0656
LSTM Sequence Model	0.0874	0.0613	0.9965	0.9916	+0.276	0.0658
XGBoost Regressor (Δ)	0.0746	0.0529	0.9972	0.9912	+0.467	0.0594
LightGBM Regressor (Δ)	0.0711	0.0495	0.9974	0.9921	+0.513	0.0587
PI-STGNN (Proposed Custom)	0.0071	0.0046	0.99998	0.9983	+0.995	0.0073




Figure 8: **Multi-Station Hydrograph Forecasts and Residual Error Distributions ($y - \hat{y}$) on Holdout Test Split.** Evaluated during the 2024 monsoon flood season (June–October 2024) across Goalundo (SW91.9R), Bhagyakul (SW93.4L), and Sureswar (SW95). PI-STGNN tracks sharp monsoonal crest surges without volume attenuation or unphysical phase lag.




Figure 9: **Quantitative Benchmark and Physical Conservation Audit.** (a) Ranked horizontal bar chart comparing Persistence Index (PI) across all baseline models and the custom PI-STGNN against naive persistence. (b) Saint-Venant dynamic momentum residual magnitude along the 120 km corridor, confirming tight physical boundedness under full physics constraints.




Figure 10: **Physics Loss Ablation Analysis.** (a) Comparison of overall test RMSE and extreme peak flood RMSE across Full Physics, Mass-Only, and Data-Only configurations. (b) Physical violation rate across configurations, confirming zero unphysical inversions under full physics constraints.

Table 11: Physics Loss Ablation Study on Out-of-Distribution Peak Flood Forecasts.

Model Configuration	Loss Weights	Mean RMSE (m)	Peak RMSE (m)	Physical Violation Rate
Full Physics PI-STGNN	$\lambda_{\text{mass}} = 0.15, \lambda_{\text{mom}} = 0.10$	0.0071	0.0073	0.00%
Mass-Only STGNN	$\lambda_{\text{mass}} = 0.15, \lambda_{\text{mom}} = 0.00$	0.0108	0.0132	0.00%
Data-Only STGNN	$\lambda_{\text{mass}} = 0.00, \lambda_{\text{mom}} = 0.00$	0.0125	0.0122	8.42%






Figure 11: **Delta Target Formulation (Δ_t) Performance Comparison.** (a) Persistence Index (PI) for Direct Level Target vs Delta Target across all 15 station-period strata. (b) Normalized error relative to mean daily river movement ($\text{RMSE}/|\Delta|$), proving that the delta formulation beats persistence in 100% of test strata.



(a) Permutation feature importance by station and lag.

Does a model fitted in one bridge period hold in another?

test-block RMSE in metres; the diagonal is the within-period score (bold)



(b) Cross-period transfer matrix (RMSE, m).

Figure 12: **Feature Attribution and Cross-Epoch Transferability.** (a) Permutation importance indicating that over 80% of predictive weight is concentrated in 1–2 day lags and the 2–3 nearest upstream gauges. (b) Cross-period transfer matrix showing generalization across BEFORE, DURING, and AFTER bridge construction epochs.

Residual diagnostics, pooled model, test block

rows: residual against fitted · normal quantile plot · autocorrelation of residuals (red = 95% band)




Figure 13: **Residual Error Diagnostics.** Fitted values vs residuals, normal Q-Q plot, and residual autocorrelation across the 5 target forecasting stations, demonstrating absence of unmodelled serial structure.

Flood-risk classification, test block, all years pooled

counts, shaded by row share




Figure 14: **Flood Risk Classification Confusion Matrices.** Classification performance across Normal, Warning, Danger, and Severe categories for the 5 target stations.

11 Usage Notes and Reuse Potential

- **Data Loading:** Load the primary panel `processed/pointwise_hydromet_7day_long.csv` with `parse_dates=['Date']`. Filter complete cases with `df[df.Window_Complete]`.
- **Target Wide Matrices:** Use `processed/wide/wide_<station>.csv` alongside feature definitions in `metadata/feature_sets.json` (`stage_only`, `hydromet`, `full_multivariate`, `forecast_safe`).
- **Recommended Modeling Formulation:** Predict the daily stage change ($\Delta = WL - WLD-1$) rather than raw stage level to eliminate piecewise-constant quantization penalties.
- **Required Baseline Standard:** Always evaluate models using the Persistence Index (PI) and Kappa Skill Score (SS_k); avoid relying solely on raw R^2 or NSE.
- **Precaution on Lead Time:** Same-day upstream features (`Pnn_WL`, `Pnn_Rain`) represent spatial nowcasting. For true 1-day forecasting, strictly use the `forecast_safe` feature subset.
- **Unsuitable Uses:** Do not perform random k-fold cross validation on raw stage levels, as serial autocorrelation will produce artificially inflated metrics.

12 Limitations and Known Biases

1. **Tarpasa (SW94) Administrative Gap:** Carries an administrative gap of 1,369 consecutive days (2016–2020), requiring its exclusion from predictor chains to preserve sample depth for downstream Sureswar.
2. **Downstream Estuarine Reach Representation:** Limited to Sureswar (31 km downstream), where localized pier afflux signals are attenuated by tidal dynamics.
3. **Cross-Period Transfer Sample Asymmetry:** Partially confounded by training sample size (DURING spans 7.5 years vs 3.5 years for BEFORE/AFTER).
4. **Historical Datum Corrections:** Gauge datum corrections across historical surveys carry residual ± 2 cm datum uncertainty.

13 Ethics, Privacy and Responsible Data Use

Table 12: Ethics, Privacy, and Responsible Data Use Disclosures.

Item	Disclosure Statement
Human participants / consent	Not applicable. Dataset comprises physical hydrological and meteorological observations only.
Privacy / de-identification	Not applicable. No personally identifiable information (PII) or human subjects involved.
Web / social-media data	Not applicable. No web scraped or social media data used.
Copyright / third-party data	Official hydrometric records procured under BWDB paid commercial invoice 2608295209 for academic research.
AI-generated / synthetic content	Zero synthetic or GenAI-generated observations were used in collection, annotation, or augmentation.

14 Data and Code Availability

Table 13: Data and Code Availability Statements.

Item	Details / URL
Repository name	https://github.com/Shakhoyat/ganges-padma-hydrometerological-dataset
Dataset DOI / persistent ID	Invoice 2608295209 (BWDB Official Archive) / GitHub Release v2.0
Direct dataset URL	https://github.com/Shakhoyat/ganges-padma-hydrometerological-dataset
Access instructions	Fully open and accessible in repository; design matrices and scripts open for evaluation
Dataset license	Creative Commons Attribution 4.0 International (CC BY 4.0)
Code / notebook URL	https://github.com/Shakhoyat/ganges-padma-hydrometerological-dataset
Code version / commit	v2.0 Release (Commit: <code>seed-20260913-final</code>)
README & environment	README.md, DATA_IN_BRIEF.md; Python 3.10+ (PyTorch, XGBoost, LightGBM, Scikit-Learn)

15 Student Contributions (CRediT-style)

Table 14: CRediT Contributor Roles and Individual Responsibilities (Matching CSE 4112 Template Table 12).

Roll	Student	Overall Contrib. (%)	Main contributions	Evidence / files owned
2107104	Md. Shakhoyat Rahman Shujon (Lead)	24%	Conceptualization, PI-STGNN Architecture Design, Multi-variable Feature Engineering, Pipeline Leadership	benchmarks/run_all_benchmarks.py, build/01_extract.py, build/04_wide.py
2107111	Mohammad Moin Uddin Moin	23%	Machine Learning Modeling, Baseline Model Tuning (LSTM, Transformer, XGBoost, LightGBM), Delta Formulation (Δ_t)	benchmarks/run_all_benchmarks.py, results/benchmark_models.csv
2107119	Md. Tariful Islam Jony	23%	Data Curation & Extraction, 7-Day Hydrometeorological Lag Structuring, BWDB Ingestion, Provenance Ledger	build/01_extract.py, build/02_registry.py, build/05_stats.py, gap_inventory.csv
2107102	Kamrul Islam	15%	Technical Validation, 3-Hourly Tidal Ground-Truth Verification, Conceptual HBV Model Implementation	results/dailyavg_validation.csv, benchmarks/run_all_benchmarks.py
2107093	Tanha Sayed Ahona	15%	Visualization & Mapping, Longitudinal Elevation Profile Generation, LaTeX Documentation, FAIR Compliance	figures/, build/10_tex.py, DATA_IN_BRIEF.md, Dataset Report

16 Acknowledgements, Funding and Competing Interests

16.1 Acknowledgements / Funding

We acknowledge the Bangladesh Water Development Board (BWDB) Hydroinformatics and Flood Forecasting Circle for providing the hydrometric data under paid invoice 2608295209.

16.2 Declaration of Competing Interests

The authors declare no competing financial or personal interests relevant to this dataset report.

17. References

- [1] Bangladesh Water Development Board (BWDB), “Hydrometric Stage, Rainfall, and Discharge Observations along the Ganges–Padma Corridor (2011–2025),” Invoice 2608295209, Dhaka, Bangladesh, 2026.
- [2] P. K. Kitanidis and R. L. Bras, “Real-time forecasting with a conceptual hydrologic model: 2. Applications and results,” *Water Resources Research*, vol. 16, no. 6, pp. 1034–1044, 1980.
- [3] J. E. Nash and J. V. Sutcliffe, “River flow forecasting through conceptual models part I: A discussion of principles,” *Journal of Hydrology*, vol. 10, no. 3, pp. 282–290, 1970.

- [4] H. V. Gupta, H. Kling, K. K. Yilmaz, and G. F. Martinez, “Decomposition of the mean squared error and NSE performance criteria: Implications for improving hydrological modelling,” *Journal of Hydrology*, vol. 377, no. 1–2, pp. 80–91, 2009.
- [5] L. Breiman, “Random Forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [6] F. Pedregosa et al., “Scikit-learn: Machine Learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [7] Flood Forecasting and Warning Centre (FFWC), “Annual Flood Report 2024,” Bangladesh Water Development Board, Dhaka, Bangladesh, 2024.
- [8] Padma Multipurpose Bridge Project Authority, “Hydraulic and Morphological Impact Monitoring Report,” Bridges Division, Ministry of Road Transport and Bridges, Dhaka, 2023.
- [9] F. Kratzert, D. Klotz, C. Brenner, K. Schulz, and M. Herrnegger, “Toward improved predictions in ungauged basins: Exploiting the power of machine learning for large-scale hydrological modelling,” *Hydrology and Earth System Sciences*, vol. 23, no. 12, pp. 5089–5110, 2019.
- [10] M. Raissi, P. Perdikaris, and G. E. Karniadakis, “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations,” *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019.
- [11] N. Addor, A. J. Newman, N. Mizukami, and M. P. Clark, “The CAMELS data set: Hydrometeorological meteorology for large-sample studies,” *Hydrology and Earth System Sciences*, vol. 21, no. 10, pp. 5293–5313, 2017.
- [12] C. Klingler, K. Schulz, and M. Herrnegger, “LamaH-CE: Large-sample data for hydrological and environmental sciences for Central Europe,” *Earth System Science Data*, vol. 13, no. 9, pp. 4529–4565, 2021.
- [13] F. Kratzert et al., “Caravan: A global community dataset for large-sample hydrology,” *Scientific Data*, vol. 10, no. 1, p. 61, 2023.

A Complete Data Dictionary and Feature Mappings

The complete data dictionary encompassing all 31 columns of the primary panel (`processed/pointwise_hydromet_7day_long.csv`) and up to 245 feature variables across wide target matrices is archived in machine-readable JSON format within `metadata/feature_sets.json`. Feature subsets are categorized into: (i) `stage_only` (pure autoregressive lags), (ii) `hydromet` (stage + rainfall + discharge), (iii) `full_multivariate` (all 31 panel covariates), and (iv) `forecast_safe` (strictly antecedent $t - 1 \dots t - 7$ lags for causal operational prediction).

B Supporting Material and Reproducibility Notes

- **Computational Environment:** Python 3.10+, PyTorch 2.1.0+cu121, LightGBM 4.6.0, XGBoost 3.4.0, Scikit-learn 1.4.0, Pandas 2.2.0, NumPy 1.26.0.
- **Random Seeds:** Pinned to 20260913 across all benchmark scripts, data splits, and model parameter initializations.
- **Hardware Execution:** Benchmark execution evaluated on an Intel Core i7 / AMD Ryzen workstation with NVIDIA RTX GPU acceleration.

C Dataset Release and FAIR-Readiness Checklist

Table 15: FAIR-Readiness and Quality Assurance Checklist (Matching CSE 4112 Template Table 13).

Checklist Item	Criterion	Yes/No	Evidence / Note
Findable	Repository and identifier provided; meaningful file names; searchable metadata.	Yes	Local repository + invoice 2608295209; structured folder hierarchy; <code>feature_sets.json</code> metadata.
Accessible	Access instructions clear; files freely downloadable in standard formats.	Yes	Standard UTF-8 CSV files directly accessible in repository with clear setup instructions.
Interoperable	Standard formats and units used; schema and codings thoroughly documented.	Yes	UTF-8 CSV, JSON, SI/metric units (mMSL, mm, m ³ /s, ppm), standard ISO dates.
Reusable	License explicit; README, data dictionary, provenance, and limitations provided.	Yes	CC BY 4.0 license; comprehensive data dictionary, limitations, and reproducible build pipeline.
Raw/Processed Split	Original raw records isolated from processed/derived tables.	Yes	Raw records stored in <code>raw/</code> ; clean panels in <code>processed/</code> and <code>processed/wide/</code> .
Versioning	Dataset version/date and change history explicitly recorded.	Yes	Version v2.0 recorded with explicit date and commit hash metadata.
Integrity	Continuous calendar reindexing; 0 duplicate dates; zero synthetic distortion.	Yes	Calendar reindexing, duplicate removal, MD5 checksum verification across all matrices.
Leakage Safety	Strict chronological partitioning; zero temporal leakage into test sets.	Yes	Chronological 80/20 train/test split within bridge epochs; out-of-bag parameter tuning.
Reproducibility	Automated end-to-end execution from raw workbooks to benchmark evaluation.	Yes	Full execution via <code>build/</code> and <code>benchmarks/</code> scripts with fixed random seed 20260913.
Ethics/Privacy	Zero human subjects; zero PII; official institutional data procurement.	Yes	Procured under official BWDB delivery; zero personal data; zero synthetic data.
Documentation	README, captions, and data dictionaries match repository schema.	Yes	<code>README.md</code> , <code>DATA_IN_BRIEF.md</code> , and LaTeX tables match all schema definitions.
Team Accountability	All five students' contributions transparently recorded.	Yes	Individual CRediT roles and contribution percentages documented for all 5 team members.